

Lift: The Language, The IR and Code Generation

Naums Mogers

Based on (Steuwer, Remmelg, Dubach, 2017)


April 2nd, 2018

University of Edinburgh

LIFT – An Intermediate Language

$$\begin{aligned}\mathbf{mapSeq}(f, \boxed{x_n \mid \cdots \mid x_2 \mid x_1}) &= \boxed{f(x_1) \mid f(x_2) \mid \cdots \mid f(x_n)} \\ \mathbf{reduceSeq}(z, f, \boxed{x_n \mid \cdots \mid x_2 \mid x_1}) &= \boxed{f(\cdots (f(f(z, x_1), x_2) \cdots), x_n)} \\ \mathbf{id}(\boxed{x_n \mid \cdots \mid x_2 \mid x_1}) &= \boxed{x_n \mid \cdots \mid x_2 \mid x_1} \\ \mathbf{iterate}^m(f, \boxed{x_n \mid \cdots \mid x_2 \mid x_1}) &= \underbrace{f(\cdots (f}_{m \text{ times}}(\boxed{x_n \mid \cdots \mid x_2 \mid x_1})))\end{aligned}$$

Data Layout Patterns



- Do not perform any computation
- Reorganize the data layout (**View**)

Data Layout Patterns

$$\begin{aligned} \text{gather}(f, \begin{array}{|c|c|c|c|} \hline x_{f(1)} & x_{f(2)} & \cdots & x_{f(n)} \\ \hline \end{array}) &= \begin{array}{|c|c|c|c|} \hline x_1 & x_2 & \cdots & x_n \\ \hline \end{array} \\ \text{scatter}(f, \begin{array}{|c|c|c|c|} \hline x_1 & x_2 & \cdots & x_n \\ \hline \end{array}) &= \begin{array}{|c|c|c|c|} \hline x_{f(1)} & x_{f(2)} & \cdots & x_{f(n)} \\ \hline \end{array} \end{aligned}$$

```
val transposeFunction = (outerSize: ArithExpr, innerSize: ArithExpr) =>
(i: ArithExpr, t: Type) => {
  val col = (i % innerSize) * outerSize
  val row = i / innerSize

  row + col
}

val Transpose = Split(N) o Gather(IndexFunction.transposeFunction(M, N)) o Join()
```

For examples of **Gather** and **Scatter** indexing functions, see
<src/main/ir/ast/package.scala>

Data Layout Patterns

$$\begin{aligned} \mathbf{zip} & \left(\begin{array}{|c|c|c|c|} \hline x_1 & x_2 & \dots & x_n \\ \hline \end{array}, \begin{array}{|c|c|c|c|} \hline y_1 & y_2 & \dots & y_n \\ \hline \end{array} \right) \\ &= \begin{array}{|c|c|c|c|} \hline (x_1, y_1) & (x_2, y_2) & \dots & (x_n, y_n) \\ \hline \end{array} \\ & \quad \mathbf{get}_i((x_1, x_2, \dots, x_n)) = x_i \end{aligned}$$

$$\begin{aligned} \mathbf{slide}(size, step, & \begin{array}{|c|c|c|c|c|c|c|c|c|c|} \hline x_1 & x_2 & \dots & \dots & \dots & \dots & \dots & \dots & \dots & x_n \\ \hline \end{array}) \\ &= \begin{array}{|c|c|c|c|c|c|c|c|c|c|} \hline x_1 & x_2 & \dots & \dots & \dots & \dots & \dots & \dots & \dots & x_n \\ \hline \end{array} \end{aligned}$$

The diagram illustrates the `slide` function. It shows a sequence of elements x_1, x_2, \dots, x_n . A window of size `size` (indicated by a double-headed arrow) is shown at the beginning of the sequence. A second window, shifted by `step` (indicated by a single-headed arrow), is shown overlapping the first. The sequence is partitioned into segments of length `size`, with the second segment starting at `step` from the beginning.

- `mapWrg`^{0,1,2}
- `mapLcl`^{0,1,2}
- `mapGlb`^{0,1,2}

- `mapWarp`
- `mapLane`
- `mapAtomWrg`
- `mapAtomLcl`

toGlobal toLocal toPrivate

```
MapWrg(MapLcl(toLocal(MapSeq(id))) $ X
```

- These primitives decouple the decision of *where* to store data from the decision of *how* the data is produced.

$$\mathbf{asVector}(\boxed{x_1 \mid x_2 \mid \dots \mid x_n}) = \overrightarrow{x_1, x_2, \dots, x_n}, \text{ } x_i \text{ is scalar}$$

$$\mathbf{asScalar}(\overrightarrow{x_1, x_2, \dots, x_n}) = \boxed{x_1 \mid x_2 \mid \dots \mid x_n}$$

$$\mathbf{mapVec}(f, \overrightarrow{x_1, x_2, \dots, x_n}) = \overrightarrow{f(x_1), f(x_2), \dots, f(x_n)}$$


- During code generation, the LIFT compiler transforms f into a vectorized form using OpenCL built-in vectorized arithmetic operations whenever possible.
 - In other cases, f is applied to each scalar in the vector.

All LIFT primitives are either:

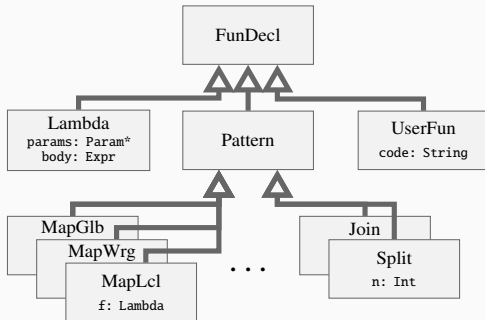
- High-level, capturing rich information about the algorithmic structure of programs
- Low-level and platform-specific (OpenCL, OpenCL for FPGAs, OpenMP, etc)

LIFT intermediate representation

Class diagram



- **Expressions** represent values and have a type associated with.



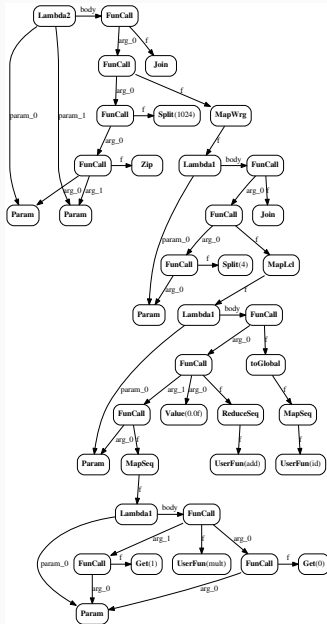
- **Function declarations** represent callable entities: lambdas, patterns and user functions.

Dot product example

```
val dotProductLift = fun(
  ArrayTypeWSWC(Float, N),
  ArrayTypeWSWC(Float, N),
  (left, right) => {
    Join() o MapWrg(
      Join() o
      MapLcl(
        toGlobal(MapSeq(id)) o
        ReduceSeq(add, 0.0f) o
        MapSeq(mult)) o
      Split(4)
    ) o Split(1024) $ Zip(left, right)
  })
```

For more dot product variations, see
<src/test/tutorial/applications/DotProduct.scala>

Corresponding AST



- Scalar types: **int**, **float**, etc
- Vector types corresponding to OpenCL types **int2**, **float4**, etc
- Tuples
 - Represented as **structs** in the generated OpenCL code
- Arrays
 - Can be nested
 - Carry information about the size and capacity of each dimension in their type
 - This information is represented by arithmetic expressions (more on this later)

LIFT compilation

Compilation stages



- Compile: [src/main/opencv/executor/Compile.scala:44](#)
 - Type checking: [src/main/ir/TypeChecker.scala:39](#)
 - Example `Pattern.checkType()`:
[src/main/opencv/ir/pattern/ReduceSeq.scala:11](#)
 - Generate: [src/main/opencv/generator/OpenCLGenerator.scala:176](#)
 - Memory address space inference:
[src/main/opencv/ir/InferOpenCLAddressSpace.scala:18](#)
 - Domain-specific range inference:
[src/main/opencv/generator/RangesAndCounts.scala:26](#)
 - Memory allocation: [src/main/ir/Type.scala:559](#)
 - Loop unrolling: [src/main/opencv/generator/ShouldUnroll.scala:50](#)
 - Barrier elimination:
[src/main/opencv/generator/BarrierElimination.scala:41](#)
 - Views (array Accesses): [src/main/ir/view/View.scala:585](#)

Memory allocation

- The naive approach would be to allocate a new output buffer for every **FunCall** AST node
- We only allocate memory to the nodes where the called function contains a **UserFun**
- The size of the memory to allocate is inferred from the array length (or the associated **View**)
- The address space is inferred from **FunCall**

Memory allocation

```

input : Lambda expression representing a program
output : Expressions annotated with address space information

inferAddressSpaceProg(lambda)
1 foreach param in lambda.params do
2   if param.type is ScalarType then param.as = PrivateMemory;
3   else param.as = GlobalMemory;
4 inferASExpr(lambda.body, null)

inferASExpr(expr, writeTo)
5 switch expr.type do
6   case Literal expr.as = PrivateMemory;
7   case Param assert (expr.as != null);
8   case FunCall
9     foreach arg in expr.args do
10       inferASExpr(arg, writeTo)
11   switch expr.f.type do
12     case is UserFun
13       if writeTo != null then expr.as = writeTo;
14       else expr.as = inferASFromArgs(expr.args);
15     case is Lambda inferASFunCall(expr.f, expr.args, writeTo);
16     case is toPrivate
17       inferASFunCall(expr.f.lambda, expr.args, PrivateMemory);
18     case is toLocal
19       inferASFunCall(expr.f.lambda, expr.args, LocalMemory);
20     case is toGlobal
21       inferASFunCall(expr.f.lambda, expr.args, GlobalMemory);
22     case is Reduce
23       inferASFunCall(expr.f.f, expr.args, expr.f.init.as);
24     case is Iterate or Map
25       inferASFunCall(expr.f.f, expr.args, writeTo);
26     otherwise do expr.as = expr.args.as;

inferASFunCall(lambda, args, writeTo)
27 foreach p in lambda.params and a in args do p.as = a.as;
28 inferASExpr(lambda.body, writeTo)
```

Algorithm 1: Recursive address space inference algorithm

Array accesses

- In LIFT IR, arrays are accessed implicitly based on the patterns
- This eliminates arbitrary memory accesses and the associated problems
- However, expressing (efficient) pattern-transformed accesses is not obvious
- ...which is where **Views** come to the rescue (but more on that later)

Barrier elimination

- We start by synchronizing after each occurrence of a parallel **Map**
- Then we remove barriers one by one in cases when it can be inferred that they are not required
 - When the data is not shared (i.e. **Split**, **Join**, **Gather** and **Scatter** are not used)
 - When the two parallel **Maps** are executed independently in separate branches of **Zip**

OpenCL code generation

- The AST is traversed recursively
- No OpenCL code is generated for the patterns that only affect **View**
- Low-level optimizations such as loop unrolling are applied to simplify the control flow using the information on *ranges* inferred from the patterns such as **mapLcl**

References

- Steuwer, Michel, Toomas Remmelg, and Christophe Dubach. "Lift: a functional data-parallel IR for high-performance GPU code generation." Code Generation and Optimization (CGO), 2017 IEEE/ACM International Symposium on. IEEE, 2017.
- Steuwer, Michel, Toomas Remmelg, and Christophe Dubach. "Matrix multiplication beyond auto-tuning: Rewrite-based GPU code generation." Compilers, Architectures, and Synthesis of Embedded Systems (CASES), 2016 International Conference on. IEEE, 2016.

<http://www.lift-project.org/>